

Vageesha Datta Ganapaneni

[✉️](mailto:vageeshadattag@gmail.com) [linkedin](#) | [Github](#) | [+1\(469\) 805-1906](#) | [🏡](#) Dublin, California

EDUCATION

• Master of Science in Computer Science

- The University of Texas at Dallas

Relevant Coursework: Machine Learning, Artificial Intelligence, Computer Vision, Natural Language Processing, Statistics in AI and ML, Web Programming Languages

TECHNICAL SKILLS

- **Languages:** Python, JavaScript, TypeScript, Rust, SQL, Bash, Linux

- **Frameworks & Libraries:** PyTorch, TensorFlow, LangChain, LangGraph, Flask, FastAPI, React, Node.js, D3.js

- **Tools & Platforms:** Docker, Kubernetes, AWS, Git, FAISS, Redis, vLLM, REST APIs

EXPERIENCE

• Handshake

Remote

- AI Fellow

Oct 2025 – Present

- Applied prompt engineering, conducted 1K+ machine learning code reviews of coding AI agents to detect rogue behaviors and integrate corrected experimentation, improving model compliance accuracy by **2%**.
- Enhancing the scientific integrity evaluation framework by designing custom metrics, contributing to the development of internal benchmarks for responsible AI governance within Handshake's MOVE research program.

• HUD

San Francisco, CA

- Research Engineer – Agentic AI Evals

Oct 2025 – Nov 2025

- Designed a distributed evaluation framework by parallelizing rollouts, vLLM inference, and fine-tuning across H100 GPU clusters, improving **evaluation throughput** by **15%** and accelerating agentic model benchmarking.
- Evaluated AI agents across multiple **codebases**, assessing their ability to interpret **pull requests (PRs)**, implement code changes, and design test cases verifying modified behavior through multi-turn reasoning.
- Refactored the evaluation harness into a modular, configuration-based system supporting plug-and-play model integration, reducing setup effort by **25%** and enabling faster experimentation for future client evaluations.
- Developed interactive **React** dashboards connected to **Python** and **Docker**-based analytics, cutting manual result inspection time by nearly **20%** and improving visibility across research and engineering teams.

• Allen Institute for AI (Ai2)

Dallas, TX

- AI Researcher

Jan 2025 – May 2025

- Applied **Optimization by Prompting (OPRO)** to enhance few-shot reasoning, refining prompt design to achieve a **12%** gain in hypothesis generation accuracy across active reasoning benchmarks.
- Developed belief-tracking and uncertainty quantification components using **entropy** and **KL divergence**, allowing early detection of reasoning drift and improving interpretability in complex inference tasks.
- Built an interactive **D3.js** platform visualizing confidence evolution across reasoning chains, streamlining analytical reviews and increasing transparency for research teams.

• IBM

India

- AI Software Engineer

Jan 2020 – May 2022

- Deployed AI-driven automation pipelines across hybrid cloud environments, improving operational throughput by **20%** and minimizing manual touchpoints in core business processes.
- Processed and structured over **10 million** data records to enable accurate training of predictive models that strengthened process reliability and decision support.
- Introduced reusable **libraries** and **REST APIs** that standardized automation routines, reducing integration timelines by **30%** and facilitating organization-wide adoption of AI tooling.
- Containerized machine learning solutions with **Docker** and orchestrated deployments using **Kubernetes**, bringing down release time from hours to minutes while ensuring consistent performance across environments.

PROJECTS

• Vision-Based Web Automation Agent: [Source Code]

- Built an autonomous web automation agent using GPT-4o Vision that executes multi-step tasks through an Observe-Think-Act loop, leveraging IoU-based filtering to deduplicate bounding boxes for accurate interactive element detection.
- Implemented intelligent action execution with visual grounding and text-matching fallbacks, achieving 98.7% accuracy in solving web automation tasks with robust loop detection and verification mechanisms.

• MoodBoard AI:Context-Aware Pin Recommendations: [Source Code]

- Built a context-aware recommendation system that generated personalized "mood boards" by combining graph embeddings of user-pin interactions with multimodal inputs, enabling richer context modeling for recommendations.
- Incorporated LLM-based embeddings of image and text metadata into the retrieval pipeline, improving recommendation quality and boosting recall by 17% in offline evaluations.